

# **Applied Containerization for Machine Learning in HPC Workshop**

## **A Hands-on Guide to Running ML Workloads in HPC Environments**

Parmanand Sinha, Computational Scientist (GIS+HPC)  
University of Chicago Research Computing Center

# Workshop Overview

- **Duration:** 2 hours (1 hour presentation, 1 hour hands-on)
- **Level:** Intermediate
- **Prerequisites:**
  - Basic programming knowledge
  - Linux CLI familiarity
  - Active RCC/cluster account
- **Repository:** [GitHub](#)

# **1. Introduction to Containerization in HPC**

# Why Containers for ML in HPC?

- **Dependency Management:**

- Package complex ML software stacks with all dependencies.
- Resolve version conflicts easily.

- **Portability:**

- Run the same environment across different HPC systems, local machines, or cloud platforms.
- "Build once, run anywhere."

# Why Containers for ML in HPC? (Cont.)

- **Reproducibility:**

- Ensure consistent results by capturing the exact software environment.
- Crucial for research and collaboration.

- **Performance:**

- Native-like performance with minimal overhead.
- Direct hardware access (e.g., GPUs, high-speed interconnects).

# Container Basics: What is a Container?

- An isolated, encapsulated user-space instance.
- Runs on a shared OS kernel but has its own:
  - Filesystem
  - Processes
  - Network interfaces (can be configured)
- Lightweight compared to Virtual Machines (VMs) as they don't require a full guest OS.



Container vs VM

Image source: [apptainer.org](https://apptainer.org)

# Container Basics: Key Components

- **Container Images:**

- Read-only templates used to create containers.
- Contain application code, libraries, dependencies, and metadata.
- Examples: Docker Hub images, SIF files (Apptainer).

- **Runtime Environments:**

- Software that runs containers (e.g., Apptainer, Docker, Charliecloud).
- Manages container lifecycle, isolation, and resource allocation.

# Containerization Workflow in HPC

```
flowchart TD
```

```
  A[Pull Container Image] --> B[Store SIF in $SCRATCH/$USER/]
```

```
  B --> C[Submit SLURM Job]
```

```
  C --> D[SLURM Allocates Compute Node]
```

```
  D --> E[Bind Data/Script Directories]
```

```
  E --> F[Run Container with Apptainer]
```

```
  F --> G[ML Training/Inference]
```



# Container Basics: Key Components (Cont.)

- **Mount Points for Data (Bind Mounts):**
  - Mechanism to make host directories/files accessible inside the container.
  - Essential for accessing datasets, scripts, and output directories.
- **Resource Allocation:**
  - Containers share host resources (CPU, memory, GPUs).
  - HPC schedulers (like SLURM) manage resource allocation for containerized jobs.

class: lead

## **2. Apptainer (formerly Singularity)**

### **Introduction**

# Apptainer & Singularity: A Quick Note

- **Apptainer** is the direct successor to **Singularity**.
- You might see `singularity` used in older documentation or as the module name on some HPC systems.
- Commands are largely interchangeable (e.g., `singularity pull` vs. `apptainer pull`).
- This workshop uses modern `apptainer` commands.

# Hands-on with Apptainer (Introduction)

*(This section introduces Apptainer. The actual hands-on will follow in the next hour.)*

# Apptainer: Getting Started

```
# Load Apptainer module (on HPC systems)  
module load apptainer  
  
# Pull a PyTorch container image from Docker Hub  
apptainer pull docker://pytorch/pytorch:latest  
# This creates a .sif file (e.g., pytorch_latest.sif)
```

- **SIF (Singularity Image Format):** Apptainer's default, optimized image format.
- **Storage Tip:** For large images, consider pulling/storing them in your scratch directory.
  - On Midway3: `$$SCRATCH/$USER/sif_files/` (where `$$SCRATCH` is `/scratch/midway3`).

# Apptainer: Basic Commands

- **Interactive Shell:**

```
# Start an interactive shell inside the container  
# --nv enables NVIDIA GPU access  
apptainer shell --nv pytorch_latest.sif
```

- **Run a Python Script:**

```
# Execute a command (e.g., a Python script) inside the container  
apptainer exec --nv pytorch_latest.sif python my_script.py
```

# Apptainer: Basic Commands (Cont.)

- **Bind Data Directories:**

```
# Run a container and mount /data on host to /data in container  
apptainer run --nv --bind /path/on/host:/path/in/container pytorch_latest.sif  
  
# Example: Mount current working directory ($PWD) to /mnt inside container  
apptainer run --nv --bind $PWD:/mnt pytorch_latest.sif
```

- `apptainer run` executes the default runscript defined in the image (if any).
- Binding `$PWD` (current working directory) is very useful for accessing your scripts and local data within the container.

# Apptainer: Key Features

- **GPU Support:** `--nv` flag for seamless NVIDIA GPU access.
- **Data Binding:** Mount host directories inside the container (`--bind` or `-B`).
  - Crucial for accessing datasets and saving results.
- **Environment Variables:** Pass configuration through the container boundary.
  - Apptainer typically inherits host environment; can be controlled.
- **MPI Support:** Run distributed workloads across nodes.
  - Often requires MPI compatibility between host and container.



class: lead

## **3. Charliecloud Overview**

# Charliecloud vs. Apptainer

Feature	Apptainer (formerly Singularity)	Charliecloud
Security Model	set-UID (optional), rootless execution	Fully unprivileged (user namespaces)
Image Format	SIF (optimized, single file)	Directory trees, SquashFS
Build Process	<code>.def</code> files, build from Docker Hub	Direct Dockerfile support
Best Use Case	Complex ML workflows, ease of use	Security-critical environments, Docker compatibility

# Charliecloud: Key Benefits

- **Minimal Attack Surface:**

- Fully unprivileged operation using user namespaces.
- Reduces security risks on shared HPC systems.

- **Docker Compatibility:**

- Direct use of Dockerfiles for building images ( `ch-convert` ).
- Easier transition for users familiar with Docker.

- **Lightweight & Simple:**

- Simpler architecture and deployment compared to some other runtimes.
- Focuses on core containerization features.

class: lead

## **4. Practical ML Container Deployment (Examples)**

# TensorFlow Example with Apptainer

```
# 1. Pull TensorFlow container (GPU version)  
apptainer pull docker://tensorflow/tensorflow:latest-gpu  
  
# 2. Run training script  
# Assumes train.py and data are accessible via bind mounts  
apptainer exec --nv tensorflow_latest-gpu.sif \  
    python /path/to/your/train.py --data /path/to/data
```

- `--nv`: Enables NVIDIA GPU access.
- Bind mount your script directory and dataset directory.

# PyTorch Example with Apptainer (Multi-GPU)

```
# Pull PyTorch container (if not already done)  
# apptainer pull docker://pytorch/pytorch:latest  
  
# Multi-GPU training using torchrun (formerly torch.distributed.launch)  
apptainer exec --nv pytorch_latest.sif \  
    torchrun --nproc_per_node=4 /path/to/your/distributed_train.py
```

- Assumes `distributed_train.py` is set up for PyTorch DDP.
- `--nproc_per_node` should match available GPUs.

# Distributed Training Topology (Multi-Node/Multi-GPU)

```
flowchart TB
    subgraph Node1 [Compute Node 1]
        GPU1[GPU 0]
        GPU2[GPU 1]
        GPU3[GPU 2]
        GPU4[GPU 3]
    end
    subgraph Node2 [Compute Node 2]
        GPU5[GPU 0]
        GPU6[GPU 1]
        GPU7[GPU 2]
        GPU8[GPU 3]
    end
    UserScript([distributed_train.py])
    UserScript -- Launches via torchrun/srun --> GPU1
    UserScript -- Launches via torchrun/srun --> GPU5
    GPU1 <--> GPU2 <--> GPU3 <--> GPU4
    GPU5 <--> GPU6 <--> GPU7 <--> GPU8
    GPU1 <--> GPU5
    GPU2 <--> GPU6
    GPU3 <--> GPU7
    GPU4 <--> GPU8
```

class: lead

## **5. SLURM Integration (Brief Overview)**



# SLURM Job Lifecycle

```
flowchart TD
```

```
  A[User Submits sbatch/srun] --> B[SLURM Scheduler]
```

```
  B --> C[Resources Allocated (Nodes/GPUs)]
```

```
  C --> D[Job Starts on Compute Node(s)]
```

```
  D --> E[Container Launched (Apptainer)]
```

```
  E --> F[ML Script Executes]
```

```
  F --> G[Job Output Written]
```

```
  G --> H[Job Completes]
```

# SLURM: Single-Node Job Example

```
#!/bin/bash
#SBATCH --job-name=ml-training
#SBATCH --gres=gpu:1          # Request 1 GPU
#SBATCH --partition=gpu       # Specify GPU partition
#SBATCH --cpus-per-task=4     # Request 4 CPUs
#SBATCH --mem=16G             # Request 16GB RAM

# Load Apptainer module
module load apptainer

# Define paths (replace with your actual paths)
CONTAINER_IMAGE=/path/to/pytorch_latest.sif
SCRIPT_DIR=/path/to/your/scripts
DATA_DIR=/path/to/your/data
OUTPUT_DIR=/path/to/your/output

# Ensure output directory exists
mkdir -p $OUTPUT_DIR

# Run the containerized job
apptainer run --nv \
  --bind $SCRIPT_DIR:/scripts \
  --bind $DATA_DIR:/data \
  --bind $OUTPUT_DIR:/output \
  $CONTAINER_IMAGE python /scripts/train.py --data_dir /data --output_dir /output
```

# SLURM: Multi-Node Distributed Training Example

```
#!/bin/bash
#SBATCH --job-name=dist-ml-train
#SBATCH --nodes=2           # Request 2 nodes
#SBATCH --ntasks-per-node=4 # 4 tasks (processes) per node
#SBATCH --gres=gpu:4        # 4 GPUs per node (total 8 GPUs)
#SBATCH --cpus-per-task=2   # 2 CPUs per task
#SBATCH --partition=gpu-multi # Example partition for multi-node GPU jobs

module load apptainer

CONTAINER_IMAGE=/path/to/your/ml_container.sif
# Script should handle distributed setup (e.g., using torchrun environment variables)

# srun will launch 'ntasks-per-node' copies of this command on each node
srun apptainer run --nv \
  --bind /path/to/data:/data \
  $CONTAINER_IMAGE \
  torchrun --nnodes=$SLURM_NNODES \
    --nproc_per_node=$SLURM_NTASKS_PER_NODE \
    --rdzv_id=$SLURM_JOB_ID \
    --rdzv_backend=c10d \
    --rdzv_endpoint=$SLURM_STEP_NODELIST:29500 \
    /path/to/your/distributed_train.py
```

class: lead

## **6. Best Practices**

# Best Practices: Container Management

- **Version Control & Tagging:**

- Tag containers with specific versions (e.g., `myimage:1.0.0`, `myimage:latest-cuda11.8`).
- Store definition files (`.def`, Dockerfiles) in version control (Git).

- **Data Management:**

- Use bind mounts for large datasets to avoid including them in images.
- Keep images small and focused on software environment.

- **Resource Allocation:**

- Match container resource needs to SLURM (or other scheduler)

# Best Practices: Performance Optimization

- **GPU Access:** Always use `--nv` (Apptainer) or equivalent for GPU workloads.
- **I/O Optimization:**
  - Bind fast storage for temporary files. On Midway3, compute nodes often have a high-throughput SSD directory at `$TEMP` (e.g., `/scratch/local/$USER/`) ideal for this.
  - For frequently accessed small files or datasets, consider staging them to such temporary storage.
  - Be mindful of I/O patterns within the container.
- **MPI Configuration:**

# Best Practices: Security Considerations

- **Unprivileged Execution:**

- Run containers without root access whenever possible.
- Apptainer runs as user by default. Charliecloud is designed for unprivileged execution.

- **Data Protection:**

- Use appropriate bind mounts; be specific about what host paths are exposed.
- Avoid overly broad mounts (e.g., binding `/`).

- **Resource Limits:**

- Rely on the HPC scheduler (SLURM) to enforce resource limits

class: lead

## **7. Additional Resources**



# Resources: Documentation

- **Apptainer Official Documentation:**

-  Apptainer Documentation

- **\*\*Charliecloud User Guide\*\*:**

-  Charliecloud User Guide

- **\*\*Workshop Repository\*\*:**

-  Workshop Repository

(Contains this presentation, examples, and hands-on materials)

# Resources: Example Repositories & Further Learning

- **ML Framework Containers (Official Hubs):**

- Docker Hub:



(e.g., `pytorch/pytorch`, `tensorflow/tensorflow`) - NVIDIA NGC:



- Docker Hub: [hub.docker.com](https://hub.docker.com/) (e.g.,

`pytorch/pytorch`, `tensorflow/tensorflow`) - NVIDIA NGC:

[ngc.nvidia.com](https://ngc.nvidia.com/) - \*\*SLURM Integration

Script\*\*  
Check your HPC center's documentation for specific

# Key Takeaway

“ **Containerization enables reproducible, portable, and efficient ML workflows in HPC environments.** ”

- Choose between **Apptainer** and **Charliecloud** based on your specific needs for:
  - Security model
  - Ease of use
  - Image format preferences
  - Integration requirements with existing Docker workflows

# **Q&A and Thank You!**

**Next: Hands-on Session!**

# References & Further Reading

- [Apptainer Documentation](#)
- [Charliecloud User Guide](#)
- [Singularity/Apptainer on RCC](#)
- [Docker Documentation](#)
- [NVIDIA NGC Containers](#)
- [PyTorch Containers](#)
- [TensorFlow Containers](#)
- [SLURM Documentation](#)
- [HPC Best Practices](#)
- [Research Computing Center UChicago](#)